A detailed historical painting depicting a large fleet of ships at sea, likely a naval battle. The ships are packed closely together, with many having multiple masts and flags. The water is depicted with various shades of blue and green, and the sky above is filled with clouds.

Scaling Fleet Clusters with gRPC

Giant Swarm GmbH

Hugo Tavares, Hector Fernandez

May 10, 2016



Outline

- ① Fleet today
- ② Why Fleet ?
- ③ Fleet gRPC
- ④ Experiment Validation Setup
- ⑤ Analysis of Results

About me

- Infrastructure Engineer @giantswarm
- Love HPC, monitoring & performance :)
- Have fun with distributed systems
- Slow runner
- Spanish who DON'T like siesta



Fleet Today

init distributed systemd

- Manage systemd units in a cluster
- Talks to systemd via dbus
- Use etcd as storage and coordination system

Fleet Today

init distributed systemd

- Manage systemd units in a cluster
- Talks to systemd via dbus
- Use etcd as storage and coordination system

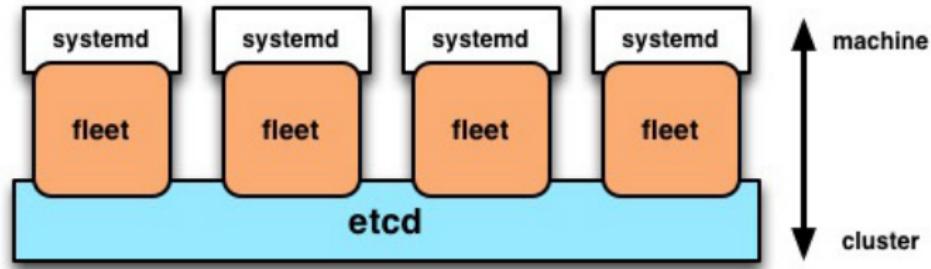


Figure: fleet overview

Fleet Today

The architecture is composed of two main building blocks:

- Engine (master)
- Agent (workers)

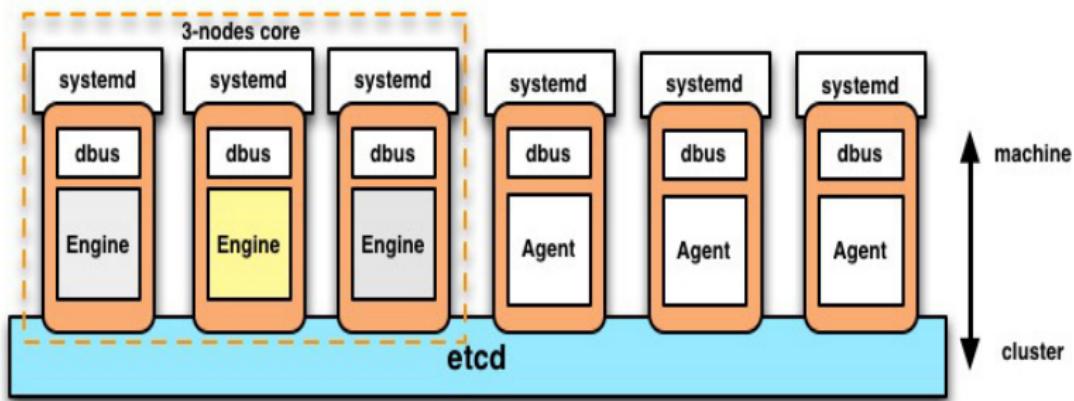
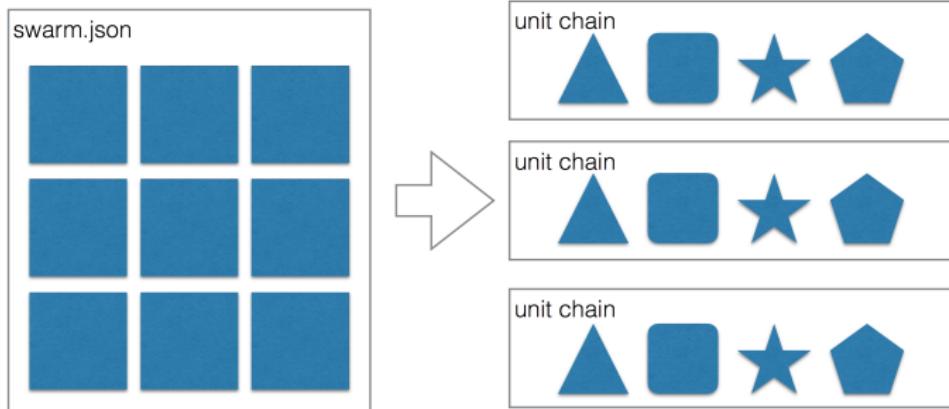


Figure: fleet: Engine and Agents

How do "we" use Fleet?

Our services are deployed as unit chains

- Service definition is done using a JSON file, 'swarm.json'





HOWEVER :/ ...

Limitations of Fleet

- Performance degradation over the time
 - actions over units taking forever
 - `etcdctl set` command can take 40ms or 1100ms
- Fleet out of sync
 - unresponsive agents
 - incoherent fleet ‘list-units’
- Limited scalability
- Excessive **etcd** dependencies: persistent and transient storage
 - *If etcd is sad. fleet is sad. We are sad.*
- Fail recovery
 - If **etcd** goes down, everything falls apart



Why Fleet ?

Why Fleet ?

Giant Swarm intensively use **fleet** as init system of our infrastructure

Why Fleet ?

Giant Swarm intensively use **fleet** as init system of our infrastructure

What do we love ?

- Robustness
- Intuitive app definition
- Distributed systemd

Why Fleet ?

How do you manage your Kubernetes systemd units?

- Fleet could help you to control Kubernetes

UNIT	MACHINE
k8s-master-api-server.service	0fc90277.../192.168.50.33
k8s-master-controller-manager.service	0fc90277.../192.168.50.33
k8s-master-scheduler.service	0fc90277.../192.168.50.33
k8s-network-setup-env.service	0fc90277.../192.168.50.33
k8s-network-setup-env.service	d077df29.../192.168.50.32
k8s-node-kubelet@348.service	d077df29.../192.168.50.32
k8s-node-kubelet@436.service	0fc90277.../192.168.50.33
k8s-node-proxy@348.service	d077df29.../192.168.50.32
k8s-node-proxy@436.service	0fc90277.../192.168.50.33

Why Fleet ?

How do you manage your Kubernetes systemd units?

- Fleet could help you to control Kubernetes

UNIT	MACHINE
k8s-master-api-server.service	0fc90277.../192.168.50.33
k8s-master-controller-manager.service	0fc90277.../192.168.50.33
k8s-master-scheduler.service	0fc90277.../192.168.50.33
k8s-network-setup-env.service	0fc90277.../192.168.50.33
k8s-network-setup-env.service	d077df29.../192.168.50.32
k8s-node-kubelet@348.service	d077df29.../192.168.50.32
k8s-node-kubelet@436.service	0fc90277.../192.168.50.33
k8s-node-proxy@348.service	d077df29.../192.168.50.32
k8s-node-proxy@436.service	0fc90277.../192.168.50.33

How do you manage other INFRA services?

- Maybe Kubernetes is not what you want to use, e.g. cleanup tasks, monitoring



What's Next ?

What's next ?

We started investigated how to:

- Reduce the etcd dependencies
- Obtain better throughput (units/sec and amount of units)
- Improve fail recovery mechanisms
- Mitigate the scalability problems

What's next ?

We started investigated how to:

- Reduce the etcd dependencies
- Obtain better throughput (units/sec and amount of units)
- Improve fail recovery mechanisms
- Mitigate the scalability problems

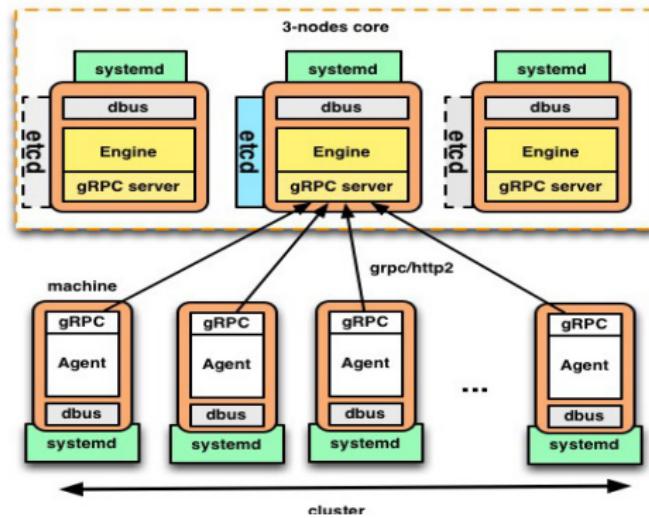
Maybe we can use gRPC instead of etcd !

Fleet gRPC

- **gRPC** is the new communication layer between: Engine & Agents
- **etcd** is *ONLY* used for *PERSISTENT* storage
- *TRANSIENT* data is stored **in-memory**

Fleet gRPC

- **gRPC** is the new communication layer between: Engine & Agents
- **etcd** is *ONLY* used for *PERSISTENT* storage
- *TRANSIENT* data is stored **in-memory**



Fleet gRPC Benefits

Fail recovery:

- Connection retry
- Resistant to network hiccups
- Fallback mechanism (e.g. firewall)
- Agent healthchecks
- If **etcd** goes DOWN, we are still operative (on-going)

Fleet gRPC Benefits

Smooth Rolling Updates

- Engine/Agents detect if there are **etcd** nodes in the cluster
 - Tested in our production INFRA (hybrid fleet cluster)

Fleet gRPC Benefits

Smooth Rolling Updates

- Engine/Agents detect if there are **etcd** nodes in the cluster
 - Tested in our production INFRA (hybrid fleet cluster)

Other benefits:

- Hybrid etcd/gRPC cluster (etcd and grpc nodes)
- New leader election algorithm
- TLS communication



Let's validate those benefits!

Experimental Validation

Configuration:

- CoreOS cluster with 3 machines as master/engine and 4 workers/agents
- systemd v219
- etcd v2.1.2
- fleet v0.11.7
- nomi v0.1.0

Tested versions:

- Scenario 1: 3k units **fleet-etcd** vs **fleet-grpc**
- Scenario 2: 20k units with **fleet-grpc**
- Scenario 3: 1k units **fleet-etcd** with **etcdv2.2** vs **etcdv3**

systemd vs fleet

How much time does it take to create 500 systemd units in 1 node?

systemd vs fleet

How much time does it take to create 500 systemd units in 1 node?

- With **fleet-grpc**: 58.8 secs
- With **fleet-etcd (current upstream version)**: 74.4 secs
- With **systemd**: 162 secs (sequential)

systemd vs fleet

How much time does it take to create 500 systemd units in 1 node?

Why ?

- ‘**systemd linking**’ operation is very EXPENSIVE

Fleet gRPC - Throughput

How do they perform to **load/start 3k units** ?

- **fleet-etcd:**

- 75 % of units started in between 1.729 - 28.14 secs
- Max: 265.88 secs

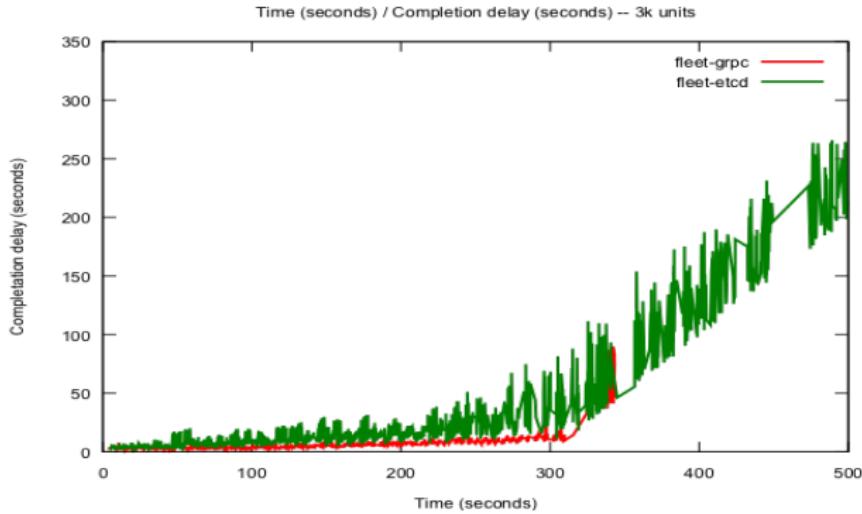
- **fleet-grpc:**

- 85 % of units started in between 1.004 - 9.907 secs
- Max: 90.0 secs

Fleet gRPC - Throughput

How do they perform to **load/start 3k units** ?

- **fleet-etcd**: 75 % of units started in between 1.729 - 28.14 secs
- **fleet-grpc**: 85 % of units started in between 1.004 - 9.907 secs



Fleet gRPC - Threshold

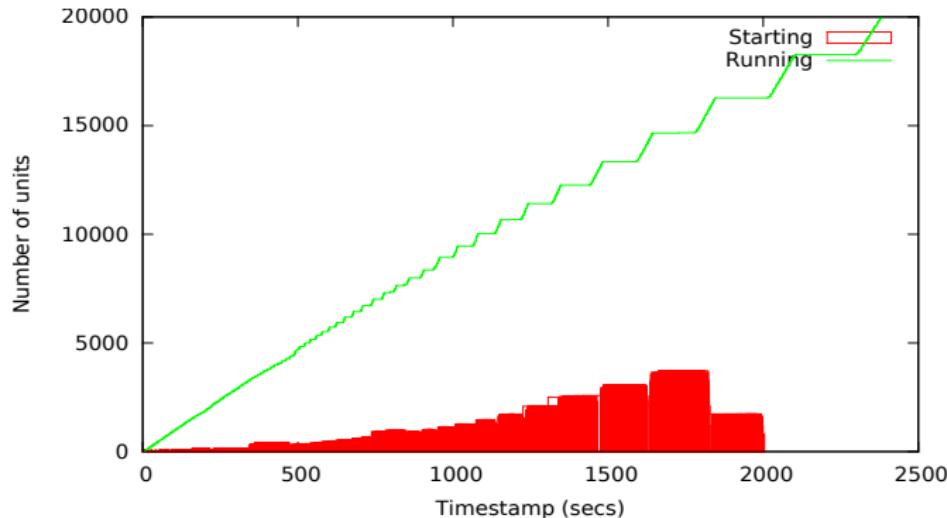
How many units can we handle in the cluster ?

- **fleet-etcd**: no more than 3k
- **fleet-grpc**: up to 20k and beyond..

Fleet gRPC - Threshold

How many units can we handle in the cluster ?

- **fleet-etcd**: no more than 3k
- **fleet-grpc**: up to 20k and beyond..



Fleet gRPC - CPU usage

CPU usage of **fleet** & **etcd** when using **fleet-etcd** ?

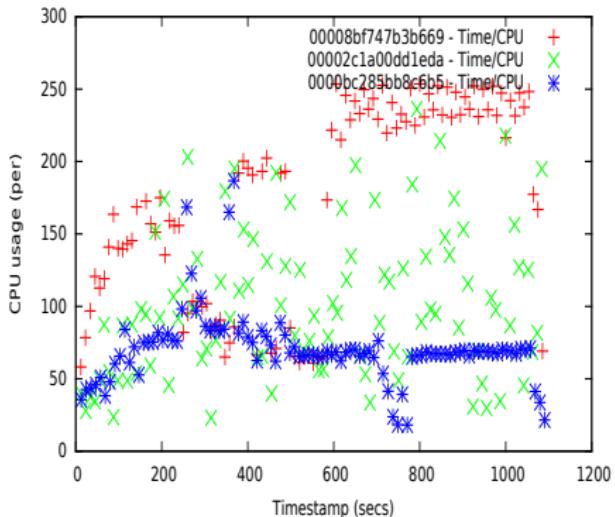
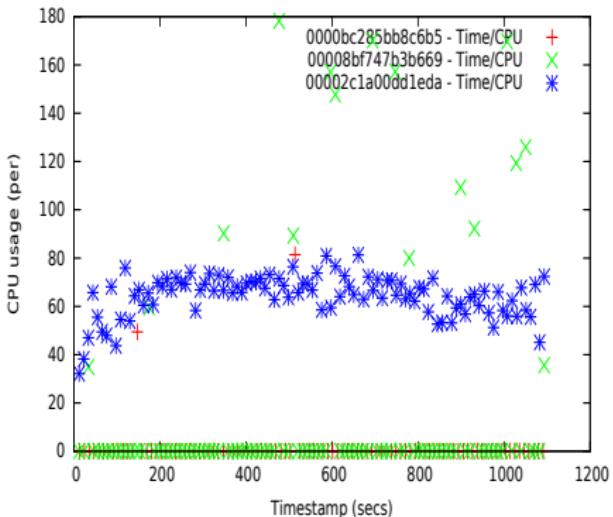


Figure: etcd / fleet – CPU usage

Fleet gRPC - CPU usage

CPU usage of **fleet** & **etcd** with the **fleet-gRPC** ?

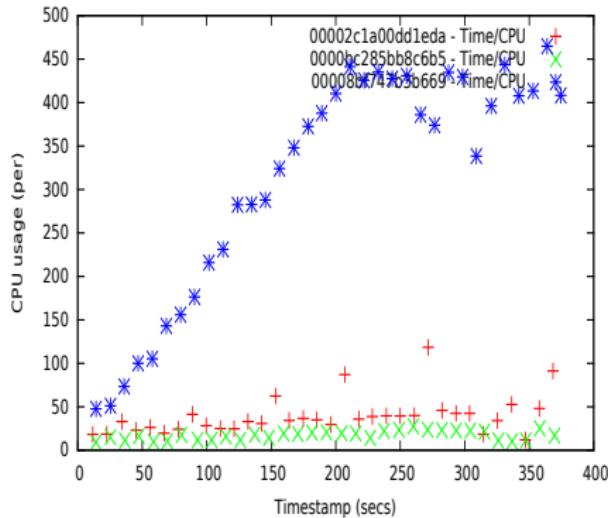
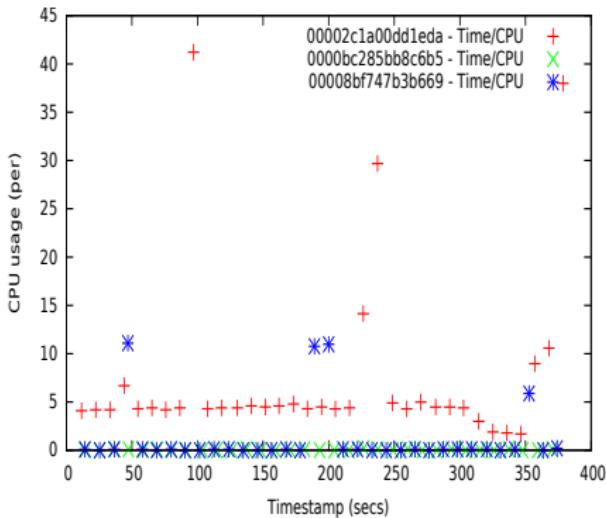


Figure: etcd / fleet – CPU usage

Fleet-gRPC - What about fleet-etcdv3 ?

Experiment:

- Evaluate the **DELAY** when starting (load/start) 1k units

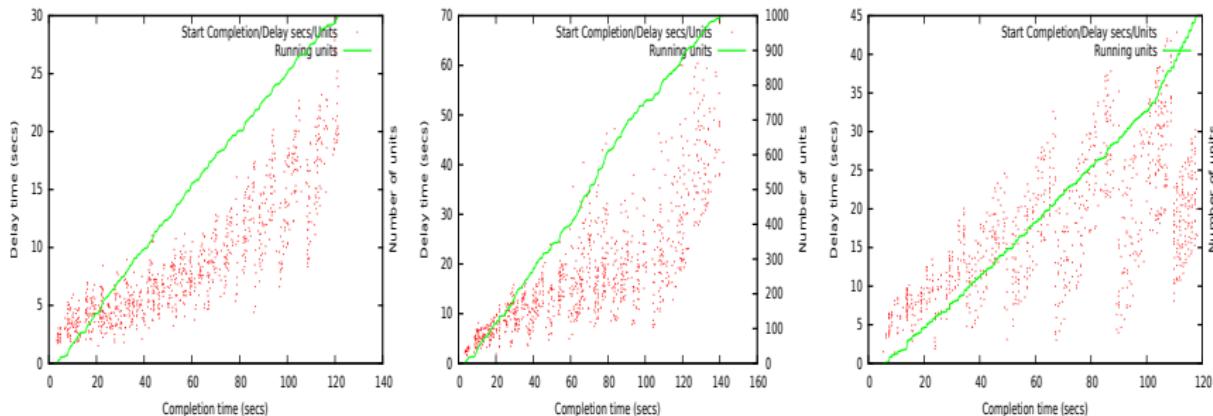


Figure: fleet-etcd v3 experimental / fleet-etcd v2.3 / fleet-grpc with etcdv2.2

Conclusion

Fleet with gRPC offers:

- Better throughput
- Less etcd friction
- Lower CPU usage in the fleet agents
- Secure communication
- Improved scalability
- No need of frequent cleanup tasks

Future Improvements

We plan to:

- Use **etcdv3**
- Add TLS to secure the communication
- Improve the unit states transmission between agents-engine, (send what it has changed)
- Change the reconciliation loop mechanism
- Look at the usage of **systemd –user** to gain in performance
 - **systemd linking** is a quite expensive operation.

Future Improvements

On-going work (PRs/Issues)

- Our fleet repo 'github.com/giantswarm/fleet'
- Our benchmarking tool 'github.com/giantswarm/nomi'

Issues

- #1426
- #1418
- #3142
- #150
- #3182
- #1393





THANKS ! Time for